

PREDICTION TASK  What is the type of task? Which entity are predictions made on? What are the possible outcomes to predict? When are outcomes observed? TASK TYPE <i>Sequence prediction (i.e., predict the next access key in a sequence).</i> PREDICTION ENTITIES <i>Data keys (i.e., unique identifiers for data items).</i> PREDICTION OUTCOMES <i>Logits representing the likelihood of each key of being accessed at each future step.</i> OUTCOMES OBSERVATION <i>At the end of the predicted future time horizon.</i>	DECISIONS  How are predictions turned into actionable recommendations or decisions for the end-user? (Mention parameters of the process / application for this.) FROM PREDICTIONS TO DECISIONS <i>Model logits and variances (over a predefined time horizon) are combined into normalized eviction scores via confidence-weighted logic (for a given confidence level). These scores are used to rank cache keys, where higher scores indicate lower eviction priority, and keys with the lowest scores—excluding protected ones—are selected for eviction.</i>	VALUE PROPOSITION  Who is the end beneficiary, and what specific pain points are addressed? How will the ML solution integrate with their workflow, and through which user interfaces? END BENEFICIARY <i>Operations/platform teams integrating the solution into a specific system (direct) and its corresponding end-users (indirect).</i> PAIN POINTS ADDRESSED 1) <i>Ineffective cache eviction strategies</i> - Rigid, rule-based policies - Poor suitability for real-world, non-stationary data access patterns - No use of predictive or probabilistic information 2) <i>Inefficient resource utilization</i> - Low cache hit rate - Suboptimal use of fast cache memory 3) <i>High operational and infrastructure costs</i> - Increased load on backing services - Higher infrastructure and operational expenses 4) <i>Degraded service performance and User Experience (UX):</i> - Increased end-user latency - Lower quality of service - Negative impact on UX - Risk of violating Service Level Objectives (SLOs)	DATA COLLECTION  How is the initial set of entities and outcomes sourced (e.g., database extracts, API pulls, manual labeling)? What strategies are in place to update data continuously while controlling cost and maintaining freshness?	DATA SOURCES  Where can we get data on entities and observed outcomes? (Mention internal and external database tables or API methods.)
IMPACT SIMULATION  What are the cost/gain values for (in)correct decisions? Which data is used to simulate pre-deployment impact? What are the criteria for deployment? Are there fairness constraints?	MAKING PREDICTIONS  Are predictions made in batch or in real time? How frequently? How much time is available for this (including featurization and decisions)? Which computational resources are used? PREDICTION TYPE AND FREQUENCY <i>Real-time and event-driven, triggered whenever the cache requires an eviction decision.</i> LATENCY CONSTRAINTS <i>The end-to-end latency is bounded to 3000 ms and is fully asynchronous, occurring outside the user request critical path.</i> COMPUTATIONAL RESOURCES <i>Server-side logic and inference using CPU-based resources within a scalable microservice architecture.</i>	ML SOLUTION IN THE WORKFLOW <i>The solution is integrated via a RESTful API within a microservice-based architecture. The cache management layer provides recent access data information, current cache state, and optional configuration settings, and receives the key(s) to evict when inserting new item(s) into a full cache. The eviction decision is fully programmatic; no user interface is required.</i>	BUILDING MODELS  How many models are needed in production? When should they be updated? How much time is available for this (including featurization and analysis)? Which computation resources are used?	FEATURES  What representations are used for entities at prediction time? What aggregations or transformations are applied to raw data sources? ENTITY REPRESENTATIONS <i>Keys are represented by three types of features:</i> - Temporal: sine and cosine of timestamp (cyclical pattern) - Local: normalized frequency and recency within lookback window - Identity: raw key mapped to a learnable embedding
	MONITORING  Which metrics and KPIs are used to track the ML solution's impact once deployed, both for end-users and for the business? How often should they be reviewed?			DATA TRANSFORMATIONS AND AGGREGATIONS <i>Raw data (timestamps and keys) is processed as follows:</i> 1) <i>Remove rows with missing values</i> 2) <i>Encode timestamps trigonometrically, transforming them into sine and cosine components</i> 3) <i>Compute local frequency, representing the normalized count of key occurrences within the lookback window</i> 4) <i>Compute local recency, representing the normalized and reversed distance from key's last occurrences within the lookback window</i> <i>Preprocessed data is then structured into input sequences of fixed length, providing the model with temporal context for predictions.</i>